

Dissertation outline for *Reasons for Persons, or the Good Successor Problem*

Kanad Chakrabarti

May 12, 2024

Abstract

Recent advances in AI may have increased the risk of human extinction or disempowerment. However, they have also made tangible our possible evolution into a machine-based species. Many may feel that a world from which our biological descendants have disappeared would be a worse one, in some irreducible or objective sense. But what about a world where certain values could be passed on to a superintelligence - values that seem to be convergently present amongst human cultures, such as respect for artistic creativity or for ecological diversity? Would this still be a terrible outcome, and if so, precisely why? The topic is controversial and confusing, yet it seems useful to set out the relevant views from philosophy, ecology and astrobiology, as well as conceptual AI. The ambition is to construct a thought experiment that asks why and how humanity's technological successor should carry the torch of intelligence forward throughout our lightcone.

1 Research question

The thesis topic is as follows¹. Imagine a future world that has some combination of recognizable biological humans, machine-human

¹It is somewhat narrower in scope than the abstract implies

cyborgs, as well as various orders of AGI.² For the avoidance of doubt, such an imagined world could be one where there are no humans or other animals but there are AGIs; conversely, it could be one where there are no AGIs – just humans and animals. Leaving aside whether such a world is stable or how it would be formed, how can we think about notions of value in such a world?³

1.1 Why is this topic worth researching?

- AI risk is a ‘wicked problem’ ([Karnofsky 2022](#)) with aspects that span technology, existential/suffering risk, public policy, and philosophy ([Bostrom 2022](#)), ([Metzinger 2021](#)), ([Althaus and Gloor 2019](#)). AI is also discussed in a politically polarised context that complicates object-level issues. ([McQuillan 2022](#))
- This project is adjacent to a number of conversations in the critical (post-)humanities. I note those adjacencies, with a

²This definition would be expanded in the dissertation e.g. “recognisable biological humans” might include humans with various implants and modest genetic modifications; AGI is being used to gesture at “advanced AI” i.e. machines at near-human capability in a broad range of tasks

³For the moment, I am slightly fuzzy on what is meant by ‘world’: whether it is just Earth or a more spacefaring context. Discussions around AGI/ASI are adjacent to space exploration and astrobiology, and particularly come up in sec. ???. For more on long-term trajectories, see ([Baum and Yampolskiy 2019](#)), and for considerations of value in respect of populations see ([Greaves 2017](#)).

view to completeness, but do not view those adjacencies as central to my project.⁴

1.2 Contribution to knowledge

2 Research method

2.1 Perspective from philosophy

To a large extent this is a philosophical project, owing to its breadth, and philosophy suggests two ways of examining the topic: a bottom-up or concrete way that starts from the world we observe today, the other is a top-down and more abstract or speculative one.⁵

⁴A number of criticisms are raised about the “project” of AI, AGI, alignment including: the critical discourse around post/trans-humanism (e.g. allegations of eugenicism, universalism, technophilia, accelerationism etc.) ([Braidotti 2013](#)) ([Manzocco 2019](#)) ([McQuillan 2022](#)) ([Gebru and Torres 2024](#)).

⁵This project institutionally sits within an art practice-based programme, hence it draws from a range of disciplines and practice, acting as a meta-discipline ([FullerM]) that hopefully generate knowledge that would be less available to any one of those disciplines individually.

2.1.1 A plurality of value

In the bottom-up approach, we start from today's world, compiling a list of preferences⁶. Some of those preferences will be convergent across societies today, but a number of those will be diverse across societies, potentially incompatible, and generally not amenable to aggregation.⁷ And then we look at a range of different future world scenarios going from the multi-species one I just described sec. 1 all the way to an empty world, or a world with perhaps one or more super-intelligences. In such a case, is there anything we can say about this grid of preferences versus worlds?

2.1.2 Nowhere/nowhen

The top-down approach is to explore a very specific question: the “view from nowhere”, and set out the history of the term, including its similar terms like *sub specie eternitatis* and “the point of view of the universe”, as well as how substantively related are they. There is

⁶Following (Russell 2019) I use preferences to encompass morality, ethics, but also aesthetic choices, rather than “values”, as the latter term sometimes might lead to confusion. A broad overview of relevant topics can be found in (Gabriel 2020)

⁷Recognising that these preferences cannot be aggregated and we're not able to say anything about humanity as a whole. See (Gabriel 2020), (Weidinger et al. 2023) and Karnofsky on Harsanyi. There may be relevant insight in social sciences, e.g. Habermas' “ideal speech situation”.

an important philosophical and historical context that led to those terms, as well as the substantial challenges to these universalist perspectives, both in philosophy as well as in a broader social and cultural context.⁸ It would be useful to also examine selected non-western literature for instance Buddhism to see if there are any analogous concepts.⁹

2.2 Evolutionary perspective

A point possibly underappreciated in the humanities is how human preferences, particularly morality and ethics, and perhaps aesthetics, seem to have been the product of a particular biological and cultural evolutionary context. They are contingent, and we must recognize that some arbitrary design of mind¹⁰, might not hold these particular preferences. Depending on one's views on moral realism, the "oughts" many of us cherish (as manifested in preferences) are, to some extent, grounded in the current and historical "is"-es of the world.

⁸Besides Nagel, Wittgenstein, and Sidgwick/Singer, Bernard Williams, Elizabeth Anscombe this should include Clare Colebrook/Donna Haraway and perhaps Negarestani; critical perspectives could perhaps start with Braidotti, Wark on Anna Tsing, Moynihan

⁹Can we say a bit more on this?

¹⁰Bracketing whether such a design is something we can get to from where we are right now

Hence it seems valuable to sketch out¹¹ what the design space of possible minds is, and set out some possible implications for the reasoning in sec. 2.1.¹²

2.3 Can decision theory help?

And then lastly, try to see if utilitarianism or decision theory or any of the more cutting-edge techniques that aren't quite in the traditional philosophy repertoire, whether those can be useful to construct a new version of the argument.

2.4 Practice: LLMs as research partners

A possible practice based engagement is to — in the knowledge that it is impossible — write A software that undertakes this project. The context of practice can be brought into it building in the tradition of eurisko or CYC or other large incomplete system building projects.

Althaus, D., and L. Gloor. 2019. “Reducing Risks of Astronomical Suffering: A Neglected Priority.” <https://longtermrisk.org/red>

¹¹Ref to sources from RWTH talk

¹²As mentioned, this project sits in an Art Department, which is more forgiving and rather encourages speculation and thought-experiments. In the interests of rigour however, I make it clear when I am veering into the more speculative.

[ucing-risks-of-astronomical-suffering-a-neglected-priority/](#).

Baum, Armstrong, S. D., and R. V. Yampolskiy. 2019. “Long-Term Trajectories of Human Civilization.” *Foresight* 21 No. 1: 53–83. <https://doi.org/10.1108/FS-04-2018-0037>.

Bostrom, N. 2022. “Propositions Concerning Digital Minds and Society.” <https://nickbostrom.com/propositions.pdf>.

Braidotti, R. 2013. *The Posthuman*. Oxford: Polity.

Gabriel, Iason. 2020. “Artificial Intelligence, Values, and Alignment.” *Minds and Machines* 30 (3): 411–37. <https://doi.org/10.1007/s11023-020-09539-2>.

Gebru, Timnit, and Émile Torres. 2024. “The TESCREAL Bundle: Eugenics and the Promise of Utopia Through Artificial General Intelligence.” *First Monday*, April. <https://doi.org/10.5210/fm.v29i4.13636>.

Greaves, Hilary. 2017. “Population Axiology.” *Philosophy Compass* 12 (11): e12442. <https://doi.org/https://doi.org/10.1111/phc3.12442>.

Karnofsky, H. 2022. “The Wicked Problem Experience.” <https://www.cold-takes.com/the-wicked-problem-experience/>.

Manzocco, R. 2019. *Transhumanism - Engineering the Human Condition*. Chichester: Springer Praxis Books.

- McQuillan, D. 2022. *Resisting AI: An Anti-Fascist Approach to Artificial Intelligence*. Bristol: Bristol University Press.
- Metzinger, T. 2021. “Artificial Suffering: An Argument for a Global Moratorium on Synthetic Phenomenology.” *Journal of Artificial Intelligence and Consciousness* 1 (8): 1–24. <https://philpapers.org/rec/METASA-4>.
- Russell, S. 2019. *Human Compatible*. New York: Penguin Press.
- Weidinger, Laura, Kevin R. McKee, Richard Everett, Saffron Huang, Tina O. Zhu, Martin J. Chadwick, Christopher Summerfield, and Iason Gabriel. 2023. “Using the Veil of Ignorance to Align AI Systems with Principles of Justice.” *Proceedings of the National Academy of Sciences of the United States of America* 120 (18): e2213709120. <https://doi.org/10.1073/pnas.2213709120>.