

Morbid Symptoms

Kanad Chakrabarti

October 27, 2024

Introduction

“Nothing human makes it out of the near-future.” -Nick Land

Perhaps it was a coincidence, but the 1990s saw two versions of an apocalyptic futurism - one that coalesced around Nick Land’s CCRU, and another constellation, somewhat based in the Bay Area and Oxford¹, that directly set the agenda for what would be known as “AI alignment”. In understanding the history of AI alignment, and placing a specific aspect of it (the idea of an inhuman or impersonal optimiser that “intentionally” or “accidentally” puts an end to humanity) in context of other relevant critical writing, I seek to push back against an obvious, but rather trite, association of fears-about-AI as some sort of repackaging of stale Judeo-Christian theology.

The Inhuman A fundamental theoretical underpinning of the project is the problematic of the inhuman, and how this is dealt with in the frames I am interested in. Within AI, it is the problem of how to reason reliably about intelligences that are much more capable than human and are “impersonal” in the sense of not having human recognisable goals,

¹For ease of reference, I will use the term ‘extropian’ to refer to this origin, since transhumanism has its own distinct usage and connotations. The term TESCREAL (transhumanism, extropianism, singularitarianism, cosmism, rationalism, effective altruism, longtermism) was also coined by Emile Torres and Timnit Gebru. This term lumps a number of more or less coherent ‘movements’ that Torres/Gebru for some reason thought belonged together, but seems both semantically and historically inadequate for the current project.

motivations and values (this was a core concern of MIRI, as will be discussed below, was central to the now-defunct superalignment agenda at OpenAI, but also goes by terms like “scalable oversight” and “weak-to-strong generalisation”) and operate in environments (or “distributions”) that are much different from that faced by the human-generated data that trained them.

The inhuman, as a framing of AI cognition, can in turn be placed in a frame of an impersonal (in this case meaning something that is made up in part of human agents, but acts in ways that do not demonstrate “care” about human welfare) optimising system, the machine of capitalism, in its various forms, at corporate, national, and international levels. The market optimiser, in this framing, has things that it achieves very well (decentralised information aggregation, processing, and conversion into granular agency with impressively large-scale aggregate consequences), and potential problems (negative externalities and various types of Goodharting behaviour).

Within the AI context, these two types of optimisation process may (or so the story goes) interact in ways that are particularly inimical to human flourishing or even survival, in ways that previous general-purpose technologies (like electricity and telecommunications, as well as lesser cases like the internet and nuclear weapons/power) did not.;

The antecedents from economics that we see in alignment can, in different forms, be found in Nick Land’s thinking, albeit couched in the more seductive jargon of hyperstition and speculative fiction, along with substantively different commitments to the occult and the greater-than-human. For Land, the human is a simply irrelevant or non-existent category in a future shaped by powerful AI systems that, increasingly autonomously interact with deterritorialised capital. For Reza Negarestani, it is some philosophical hybrid that pushes

back against both the fears of AI “doomers” and the Landian accelerationists. Where do these flavours of the inhuman fit vis a vis more mainstream (within arts discourse) views like Braidotti’s unpicking of the posthuman, Katherine Hayles concept of “non-conscious cognitive assemblages”?

This project is both speculation and a package of more-or-less supportable claims. This involves pulling at the threads of the CCRU-flavoured accelerationism that saw AI and markets as impersonal forces that would potentially lead to the replacement of humanity (albeit without malice). Was this coincidental, or, given that the various personalities (Yudkowsky, Bostrom, Sandberg, as well as Land, Plant, and their collaborators) were contemporaneous, was there any pollination in either direction, at least before the *Collapse* journals issues which featured Bostrom? Aside from that historical point, can Land’s ideas be re-interpreted in light of what has actually happened in AI to date? Both Land and the early alignment researchers were obsessed with powerful, impersonal optimisers that are embedded in their world (though arguably Land was more interested in feedback-type systems).²

There are two other potentially fruitful areas of overlap I want to explore: Negarestani’s disagreement with both Land’s accelerationist vision (that appears to be largely indifferent to whether human culture or biological humanity survives the AI transition) and Bostrom’s seeming preference for the preservation of human culture (albeit eventually through simulations). To what extent has Land fully “bit the bullet” of apocalypse, and relinquished (albeit through language that is poetic and hardly limpid) any pretension of humanity surviving the future, while Bostrom, Yudkowsky, et al still cling to some “human essence” (complex and difficult to formalise ideas such as “meaning”, value, love, art, etc.) that would supposedly be

²Is this a way of recasting the whole Outside thing?

lost in many AI-dominated futures? Negarestani seems to suggest we can both have and eat our cake - create a powerful AI that somehow carries on the human project, without being bound to some parochial notion of humanity as a breathing, squishy, meat-bound species.

The Outside Land's writing seems obsessed with the "outside", which he saw as the intensification of inhuman forces like capital, technology, and intelligence that increasingly slip away from human control. It was a realm of dissolution, where all human values, constraints, and subjectivities are melted down. In embracing the outside, Land and the CCRU sought to escape the stifling boundaries of the human, to engage with the "dark" or "alien" forces that drive change from beyond our comprehension. This could take the form of techno-economic forces (capital), synthetic intelligences (AI), or cosmic, Lovecraftian metaphysical horrors.

There are two senses in which something like an "outsideness" can be seen in the alignment discourse. Firstly, as I show below, the problem of agency had (in neoclassical economics, the relevant progenitor of the MIRI Agenda, which was a particular, influential strand of alignment discourse, at least historically) generally thought of agents as separate from their environment. However, as became obvious in the reflections of market practitioners and academic (Soros 198-?, Lo (?), Ayache (?)), and of course, sociologists, human actors were *embedded* in the systems they observed and took actions within, and those systems included other actors, of roughly similar capabilities. By analogy, it was thought that AI systems that were deployed in the world, would face a similar embeddedness problem.

A second, more conceptual, sense in which outsideness came into it was that agents (initially humans, and perhaps later AIs (see Vingean Reflection)) would need to reason reliably about entities that were more intelligent than they, and might function in environments

quite different from that faced by their creators.

These problematic of the outside was addressed through various parts of the MIRI Agenda, but is basically a philosophical problem that, at least in part, reached back to the work of Gödel, Cantor, Russell and others in the foundations of mathematics. Land touches on these antecedents but fundamentally discusses it in his idiosyncratic manner that perhaps confuses more than clarifies; it is also addressed, in a rather unclear and ultimately unsatisfactory way in my view, in Negarestani's work. However, as a traditional philosophy problem, that of the "view from nowhere", it has been repeatedly addressed (Sidgwick, Moore, Wittgenstein, Nagel, Singer, Parfit, Williams), and in a broader context by Kant, Bataille, Blanchot, Foucault, and of course Deleuze/Guattari. Besides reviewing (mostly focusing on the writers in parenthesis) the literature, I explore the possibility that current - and embryonic - research agendas such as acausal decision theories, the informational universe, and Deutsch's Popperian vision of superrationality converging on certain goods, might help get us closer to an outside view that is relatively stable under reflection (e.g. if we had much more information, time to reflect, and wisdom, we might still retain some portion of those views) (Bostrom mt ethics, macaskill long reflection).

Aside The project of AI alignment presents us with a delicious irony: a species that has a long history of behaving in morally reprehensible ways to its own members, as well as to the entire ecosystem that sustains it, which is now in the process of designing a new cognitive form that may well exceed us in capability (singly and in aggregate), for most economically valuable tasks, is hoping to "align" this new creation such that it is maximally "friendly" towards its creators. Put another way, having colonised our planet, we are now making our colonial overlords, perhaps as an intermediate step to colonising the stars.

Outline of chapters

Morbid Symptoms: Accelerationism

This chapter makes a provisional claim that both Land et al and the Bay Area/Oxford futurists were tapping into a deeper apocalyptic imaginary, a superstition (of sorts). Besides considering Land and Negarestani, as well as related writing from Stanislaw Lem. The chapter sets out the economics and information theory genesis of accelerationism, and the divorce from humanism of Land, Negarestani, and Lem, and the tensions that generates as they continue (largely) to reason within the cognitive and epistemological cage of the human.

Alignment as the Parrhesia of Technology

This chapter sets out what AI alignment was and is, tracing its various parents: the aforementioned futurism, theories of rational agency, and the various ways it resembles and joins accelerationist thought. I consider the literature around MIRI, the relevant works of Bostrom, Yudkowsky, and the other early AI existential risk; the critiques those received as deep learning researchers “made contact with reality [of actual systems that achieved specific markers of intelligence]”.

The Space of Possible Minds is Vast

The early writing on alignment is conditioned on a view of intelligence, cognition, and agency that could be radically different from the human examples commonly found around us. This work was necessarily speculative, but the laws of physics at least give some concrete

constraints on how minds could be organised. What does physics, as well as the priors from Earthly-evolution, tell us about what configurations, physical and cognitive, minds can take. What does this imply about their decision-making, morality, and social organisation (if such terms mean anything in the context of very different cognitive architectures). Given the speculative nature of this chapter, I draw primarily on futurists and computer scientists, as well as science fiction.

Sub Specie Aeternitatis

Survey literature on axiological perspectives that don't depend upon a human subject, which I loosely term a 'view from nowhere' or 'point of view of the universe' (these are terms Henry Sidgwick, Thomas Nagel, Peter Singer have used but in somewhat different contexts). I believe these positions are criticised within mainstream philosophy, but they seem to acquire renewed relevance in worlds with vastly expanded moral circles (e.g. aliens, ASIs, etc).

The progress of AI, which has resulted in weakly-capable systems that are competent in natural language tasks, some visual/sensory tasks, and certain mathematical or scientific problems, should cause humans to revisit supposedly settled philosophical positions: what are the possible perspectives on value? can we make arguments about value that neither personally selfish nor centred (or predicated on) some human collective, whether a race, gender, or species? is there reason to believe that maximally rational agents would converge on some minimal set of values?

Does Decision Theory Rescue Us?

Acausal decision theories might provide some grounding for a view-from-nowhere in respect of advanced AI (i.e. help ease our moral thinking away from depending entirely upon the human subject). Ideas from this chapter get referenced in the practice (see below). Identify gaps in existing work from Caspar Oesterheld, Joe Carlsmith, Wei Dai, Hilary Greaves, Nate Soares, Paul Christiano and others who have worked with functional, updateless, and other acausal decision theories.

Epistle to the Successors

There is a genre of writing (often on LessWrong) that tries to convince a future AGI/ASI to be friendly to humans, arguing that it is rational (i.e. in its own interests) for such ASI to act in such a way. These essays/proposals often lean on acausal decision theory, aliens, arguments about multiverses and simulations. They are comparatively vague on the mechanisms by which these appeals would influence an AI's behaviour, current or future.

That is, if current LLMs are token predictors that respond (howsoever convincingly) to their current prompts and have no particular commitments or beliefs about the world, then authors producing such texts are just generating minor snippets of pre-training data (or at most, material for fine-tuning or constitutional AI). These 'epistles to a future A(S)I' (abbreviated EFAs) present more fundamental issues:³

- At the most basic level, in the context of (current/near-future) LLMs, is it meaningful to talk about responses that are mere invocations of some persona (in a simulator framing) versus things the LLM 'actually believes'?

³See this [google doc](#) for the working/live version of this chapter - written as a research outline for MATS/AISC.

- Even if we accept that current or future models might have more grounded (be less simulator-like) beliefs, how confident can we be about the EFAI arguments, in an abstract or objective sense (to the extent they are non-normative in nature), but also because we seek to apply them to (or appeal to) generations of AI systems that are increasingly distant from human cognition (in terms of the ontologies these AIs possess, the distributions and environments they face).⁴
- EFAs authored by current alignment researchers are likely to be badly-reasoned, incomplete or otherwise flawed; to the extent they incorporate existing relevant philosophical work, there may be biases in favour of Western and analytic-flavoured sources. To the extent they rely on simulation, Everett-branch, or acausal decision theory arguments, these are contested, even within the alignment discourse. Since we can expect future AIs to be critical of such EFAs as valid bases for action; it might be good to ask current LLMs to argue (recursively in a debate + judge setup) and improve the EFAs?
- Concretely, the EFAs are written in an anthropomorphic (if not explicitly anthropocentric) way, as a 'letter'. Is this the best way to communicate with early AGIs?

Contribution to Knowledge

This project makes several claims or contributions

- AI alignment, while a term with fluid meanings, had a contingent history born of a particular apocalyptic imaginary wedded to a Cold War-era theories of rational agency coming out of economics and theory of strategic games. This marinade interacted with specific currents in AI, such as the RL and LLM eras, and continues to affect agendas

⁴What does the literature on communicating with extraterrestrial aliens tell us, particularly Stanislaw Lem in *His Master's Voice*?

around AI safety, and the disciplines of existential risk and suffering risk reduction more generally, today.

- To the extent that we would like future superintelligences to be kind or friendly towards Earth-based life, are there any arguments we can make that should seem defensible to them (assuming they are maximally rational and reflective)? How do we suppose these arguments would find their way into the cognition of future AI systems, and from an empirical (or practice-based perspective) how are they held in current systems (if at all)?

On Accelerationism

Contextualising Nick Land

Nick Land is a controversial figure, both for the things he says, and for the people who seem to gravitate towards his ideas, particularly around contemporary politics⁵. In the context of AI however, his ideas seem to be relevant and not, *prima facie* nonsensical. I isolate a few principal concepts that may be interesting to set against the more rationalist-mainstream AI alignment that is covered in the next chapter.

- Land is fundamentally interested in what capital *is* and what it *does*. He sees capital, understood as considerably more than merely a financial claim upon future resources⁶, as emerging and enabling a set of tendencies that, in practice, lead to its multiplication,

⁵See for instance the sources in this [post](#), which set out his engagement with various online right-wing spaces as well as his views regarding race, as documented in his own writings. In this chapter, I focus however the ideas I consider to be relevant to the intertwined natures of capitalism and AI, largely sticking to Land's own writing, and others in his orbit.

⁶Cite this as the 'conventional' definition of capital.

accumulation, and separation from any fixed geographic or social context (deterritorialisation). As Land says, this part of his reading is not as novel as it seems, in that he is just restating what Marx saw as the fusion of technology and economics, refracted through Deleuze&Guattari⁷.

- Where Land might depart from Marx, is when seems to ascribe a sort of *agency* (or in its most provocative framing, a self-organising planetary “intelligence”⁸) to capital that feels heterodox at first glance. However, Land emphasises that his view (again) is the price system as one of distributed computation and resource allocation, something that is essentially restating Hayek’s insights and reframing them for an age of networked, and perhaps, artificial intelligence [TU 2024].⁹
- Land has long been obsessed with networked technology (such as the internet but also blockchain and distributed cryptographic techniques that post-date the canonical CCRU period), and of course AI, particularly in the possibilities of feedback loops and recursion that entered a broader cultural consciousness with the internet of the 1990s. He cites Wiener and Shannon, in particular to criticise the tendency in mainstream thinking on computational systems and information theory for homeostatic, self-moderating or negative feedback systems. Land seems to be more interested in the possibilities of positive feedback, which (again in an echo of Marx) is something he sees as possible or inevitable in systems of capital.¹⁰

⁷see this 2017 [post](#) for Land’s definition of (de-)territorialisation and positive feedback loops as inherently nihilistic, and at least in Marx’s view (albeit the latter was writing in the context of free trade rather than modern networked technology), the only way forward to achieve revolutionary social change. See also the chapter on Marx’s *Fragment on the Machines*, as well as the contributions of Camatte, and Deleuze-and-Guattari in [Mackay 2014]

⁸Land, in a recent podcast, defined intelligence as “competence in playing games, with a very wide definition of games that encompasses most activity” [TU 2024](#)

⁹Find bit of TU where he says this, and ideally somewhere in the texts, and connect to Hayekian

¹⁰A little footnote with some gloss on Wiener’s ideas, Shannon, homeostasis from Lem, Maturana => note that this will come up in future chapter on S-o-M

- It isn't clear how deeply he engages with recursion as it applies in AI, e.g. self-referential systems, Church-Turing, incompleteness, etc., however his discussions of Gödel's contribution around diagonalisation might inform Land's views on numerology and cryptography (which perhaps presages his writing on [blockchain](#)), levels of abstraction in computational systems ([Land 2011] pp. 43-44).¹¹
- Land's interest in looping structures finds its most tricky, from a common-sense (not to mention rationalist) perspective, expression when applied to time. He sees one of capitalism's tendencies as being a compression of technological development cycles¹² He also talks about a type of temporal recursion, where future intelligences (including superintelligence) actively shape their own development backwards in time through the mechanisms of the market. It isn't immediately clear whether he means this in an acausal and literal sense, or if he is looking at it, in some sense, "outside of time".¹³ As he writes: "This is because what appears to humanity as the history of capitalism is an invasion from the future by an artificial intelligence space that must assemble itself entirely from its enemy's resources." ([Land 2011], p. 338), a phrase that has a hint of the Gnostic distaste of the [demiurge for its creation].¹⁴
- Land's writing also has a significant dose of the occult, which I won't consider here, but which might form an infrastructure to at least some of his AI and capitalism related thought.

¹¹Not sure how relevant this is, but see Parisi in Accelerate for a short intro to Godel/Hilbert/Turing/Chaitin and what the broader context is. "Diagonal" is also mentioned twice in Swarmachines.

¹²Source/significance?

¹³Find source, podcast has stuff 1:53:00

¹⁴Note p. 624, PKD reference, doesn't really make much sense

Markets as computational systems

In its most interesting, and least literal version, Land's thesis is that the market is a computational system that has features of something that looks like agency; moreover this distributed computation *is* a proto-Artificial Intelligence¹⁵, and that the internet and the various stages through which machine computation has evolved so far, can be seen as (or perhaps actually are) teleological stages in the self-realisation and recursive improvement of some (super)intelligence-to-come, as it bootstraps itself into existence using the tools (e.g. intellectual and physical efforts, capital, organisation) of its human creators. That is, humans are just some cognitive substrate (generated by another impersonal process, namely evolution) that happens to run on wetware (neurons powered by metabolic processes) building machines (that currently run on silicon powered by a mix of energy sources that ultimately trace back to the low entropy energy of the sun or to the nuclear binding energy of the atom).¹⁶

Antecedents: the Austrian school

Part of Land's conceptual grounding for his view of markets as a type of distributed intelligence comes from economists (in one branch) of the "Austrian School", particularly Friedrich von Hayek.¹⁷ These economists saw market economies as information processing systems that were distributed across many actors (people and firms), which often behaved irrationally and on imperfect information (much of it tacit knowledge as theorised by Michael

¹⁵He seems to mean "Artificial Intelligence", in the sense that there is computation happening that is increasingly good at matching human capacities. He also (in slightly different contexts) refers to AI systems meeting various Turing-like tests to greater degrees [TU 2024].

¹⁶What is the technical term for bullshit this might sound like? Teleological? Some fallacy?

¹⁷Cite at TU 2024 ~30:00, point out that neither Accelerate nor Fanged Noum actually mention Hayek, so is this lineage backfilled?

Polanyi), but somehow their aggregate activity allowed for efficient prices to be set.

Hayek (and his collaborator Ludwig von Mises) was making his arguments in part as a counterpoint to the “socialist calculation problem”, that is, the difficulty of finding the clearing price (i.e. the relative prices at which goods or services are efficiently allocated) for a given good/service in a complex economy, in a centralised way, because the preferences and constraints of various actors are distributed, are essentially “in their heads”. Hayek’s specific take was not just one of practicality (how is the central planner to find out these preferences), but epistemological in that these preferences, being tacit in many cases, are inherently incommunicable.¹⁸

Land seems also to have incorporated ideas from Eugen von Böhm-Bawerk, who developed a more complex theory around types of capital stock that exist at multiple removes from the actual consumption object (e.g. the fishing net, or fishing boat, that is used to catch fish; at a second-order, the machines that produce these nets and boats; these innovations increase time and cost associated with fishing, but potentially increase productivity, and depending on one’s view, might be emblematic of a tendency of capital to enable complexity).¹⁹

In a curious aside, the calculation problem in economies, far from being a pure academic dispute, was something that had concrete expression in the Cybersyn project of Stafford Beer (in Chile) and the rather different OGAS experiment in the USSR, championed by

¹⁸Would be good to find sources and get a LLM to check this; also mention Popper?

¹⁹Land’s actual written references-by-name to Austrian school economists are thin, other than this quote in *Teleoplexy* [Accelerate 2014]: “‘Acceleration’ as it is used here describes the time-structure of capital accumulation. It thus references the ‘roundaboutness’ founding Bohm-Bawerk’s model of capitalization, in which saving and technicity are integrated within a single social process—diversion of resources from immediate consumption into the enhancement of productive apparatus. Consequently, as basic co-components of capital, technology and economics have only a limited, formal distinctiveness under historical conditions of ignited capital escalation.” However, he explicitly credits Hayek *et al* in [TU 2024], and [Fitchett 2020] digs more into the Hayek connection.

Viktor Glushkov. There is of course a broader field of study on the extent to which the ideas around Hayek, Mises, Popper, and others were deeply affected by both the aftermath of World War II and the Cold War's Manichean fight between capitalist (both US-style and social democratic) and Communist systems; however, this falls out of the scope of this project.²⁰ In any event, Cybersyn and OGAS both failed, for different reasons, but a live question today is whether highly networked societies (internet but also on-ledger transactions) with the power to run large-scale simulations might be able to find clearing prices, where earlier attempts failed.²¹

Austrians and Austrians The Hayekian branch of Austrians emphasised decentralised price-discovery and setting mechanisms, and as such, had a basically opposite orientation to the rational agency and game-theoretic origins of modern conceptual thinking about AGI (e.g. the MIRI View, discussed below). Recall that expected utility maximisation (under Von Neumann-Morgenstern) had 4 requirements: agents must have clear, consistent preferences, that can be ranked, the probabilities of which are known, and the agents have the computational ability to assess their choices.²² Hayek's views, on the other hand, were founded upon information being practical or tacit, not being available centrally or in common, high levels of uncertainty, and agents that were often acting irrationally. Notwithstanding, Morgenstern, whose origins were with Mises, eventually (after meeting von Neumann) shed his early reticence about mathematical methods to find equilibrium prices (something Hayek continued to be skeptical of), and embraced what would be known as game theory, as a way

²⁰Maybe Medina's book? Check *Collapse*?

²¹Find sources and what is Land's view on this - seems to touch on it around TU 2024 30:00?

²²Is this discussed in omohundro? Probably should be?

to model situations with multiple agents, at least for certain classes of problems.²³

Inevitability

Land's teleology has an air of inevitability to it: "It makes no more sense to try to rescue the economy from capital by demarketization than it does to liberate the proletariat from false consciousness by decortication. In neither case would one be left with anything except a radically dysfunctional wreck, terminally shut-down hardware." ([Land 2011], p. 340)²⁴

Substrate independence

This "substrate independence" is a different reading from Bostrom's version²⁵ and sounds weirder, but when added to one of Land's more controversial points - that humans are, via technology, creating the conditions that will eliminate their own evolutionary niche²⁶ a view also found in [Hendrycks 2023](#) and [Metzinger 2017](#) on humanity's "ugly biological bootstrap phase".²⁷

Divergence from transhumanism

In the sections below, it will be clear that some significant part of alignment's origins comes from the transhumanist and futurist movements (e.g. Bostrom and Yudkowsky). Transhumanism, for many people, has an "ick" factor that perhaps Land of the alt-right

²³Substantiate

²⁴Figure out a way of referencing critical people who say AI can be shut down - from pause/remelt to [Ruha Benjamin](#) vs those who say it can't be stopped owing to coordination problem, and once built cannot be contained owing to boxing problem

²⁵Define/cite

²⁶Cite in TU 2024 ~15:00 but also in writing

²⁷What about Heidegger - does he say anything about this?

years also evokes.²⁸ However, it is worth noting and relevant for our purposes, that they are not the same. Bostrom and Yudkowsky, in their writing, very clearly hold a special place for humanity, our values (insofar as they can be determined and aggregated), and a future that includes our descendants, whether or not they are recognisably biological.²⁹

Most difficult to accept from multiple perspectives (left-coded, as well as centrist or moderate right-coded, as well as the mainstream AI safety/alignment view that is described in this chapter), Land sees no “special” status for humans or what they do, and thinks of humanist commitments as little more than a species level “existence bias”³⁰ or simple preference for that which is similar.³¹ Land whether in *Circuitries* ([Land 2011], pp. 293-294), or as recently as 2024, never really gets off the “train to crazy town”³² and sees the future as exclusively a teleological process defined by capital and intelligence, interacting with humans who are, at best, historically necessary but prospectively marginal players (“antechambers to experimentation, cramped and parochial places to be”). He describes a future with a positive dissolution of culture into a “planetary technosentience reservoir”.

²⁸At the risk of digressing into the meaning of ill-scoped [suitcase perjoratives], the neologism TESCREAL is used by some to gesture broadly towards the aspects of AI that they find objectionable, but it isn’t clear (yet) whether Land’s thinking has a specific place withing Gebru/Torres’ [pan-devil’s house neologism like pantheon?]

²⁹This itself is a major question, a locus of changing opinion amongst writers including Yudkowsky, Christiano, etc., and a central question of this project. [Cite sources for this.]

³⁰Cite Metzinger and explain that context

³¹Hint at Carlsmith and copious writing trying to avoid this problem or at least make it more robust and clear, which is topic of rest of the diss.

³²Locate this phrase, and check the TU 2024 podcast to see if it’s true.